



Natural Language Processing

第三周 卷积神经网络

庞彦

yanpang@gzhu.edu.cn

Overview



CONTENTS

01

卷积神经网络

02

生成对抗网络




01

Convolutional Neural Networks

卷积神经网络


Spring 2023

Convolutional Neural Networks




Left: A regular 3-layer Neural Network. Right: A ConvNet arranges its neurons in three dimensions (width, height, depth), as visualized in one of the layers. Every layer of a ConvNet transforms the 3D input volume to a 3D output volume of neuron activations. In this example, the red input layer holds the image, so its width and height would be the dimensions of the image, and the depth would be 3 (Red, Green, Blue channels).

Features on each layer



Color Image: 3 channels



Pixel of an RGB image are formed from the corresponding pixel of the three component images

Color Image: 3 channels



Convolution Layer: The Kernel



1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		


Convolved
Feature

Operation	Filter	Convolved Image
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Edge detection	$\begin{bmatrix} 1 & 0 & -1 \\ 0 & 0 & 0 \\ -1 & 0 & 1 \end{bmatrix}$	
	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
	$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Box blur (normalized)	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	
Gaussian blur (approximation)	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	

Convolution Layer: The Kernel



Colorful image
彩色图



1	-1	-1
-1	1	-1
-1	-1	1

Filter 1


-1	1	-1
-1	1	-1
-1	1	-1

Filter 2



1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

Convolution Layer: The Kernel



0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

Input Channel #1 (Red)

-1	-1	1
0	1	-1
0	1	1

Kernel Channel #1

308

+

0	0	0	0	0	0	...
0	167	166	167	169	169	...
0	164	165	168	170	170	...
0	160	162	166	169	170	...
0	156	156	159	163	168	...
0	155	153	153	158	168	...
...

Input Channel #2 (Green)

1	0	0
1	-1	-1
1	0	-1

Kernel Channel #2

-498

0	0	0	0	0	0	...
0	163	162	163	165	165	...
0	160	161	164	166	166	...
0	156	158	162	165	166	...
0	155	155	158	162	167	...
0	154	152	152	157	167	...
...

Input Channel #3 (Blue)

0	1	1
0	1	0
1	-1	1

Kernel Channel #3

164

+


+ 1 = -25

Bias = 1


Output

-25				...
				...
				...
				...
				...
...

Convolution Layer: Padding and Stride



padding = 1
stride = 1



padding = 1
stride = 2

Pooling Layer



3.0	3.0	3.0
3.0	3.0	3.0
3.0	2.0	3.0

3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1

12	20	30	0
8	12	2	0
34	70	37	4
112	100	25	12

max pooling

20	30
112	37

average pooling

13	8
79	20


Features



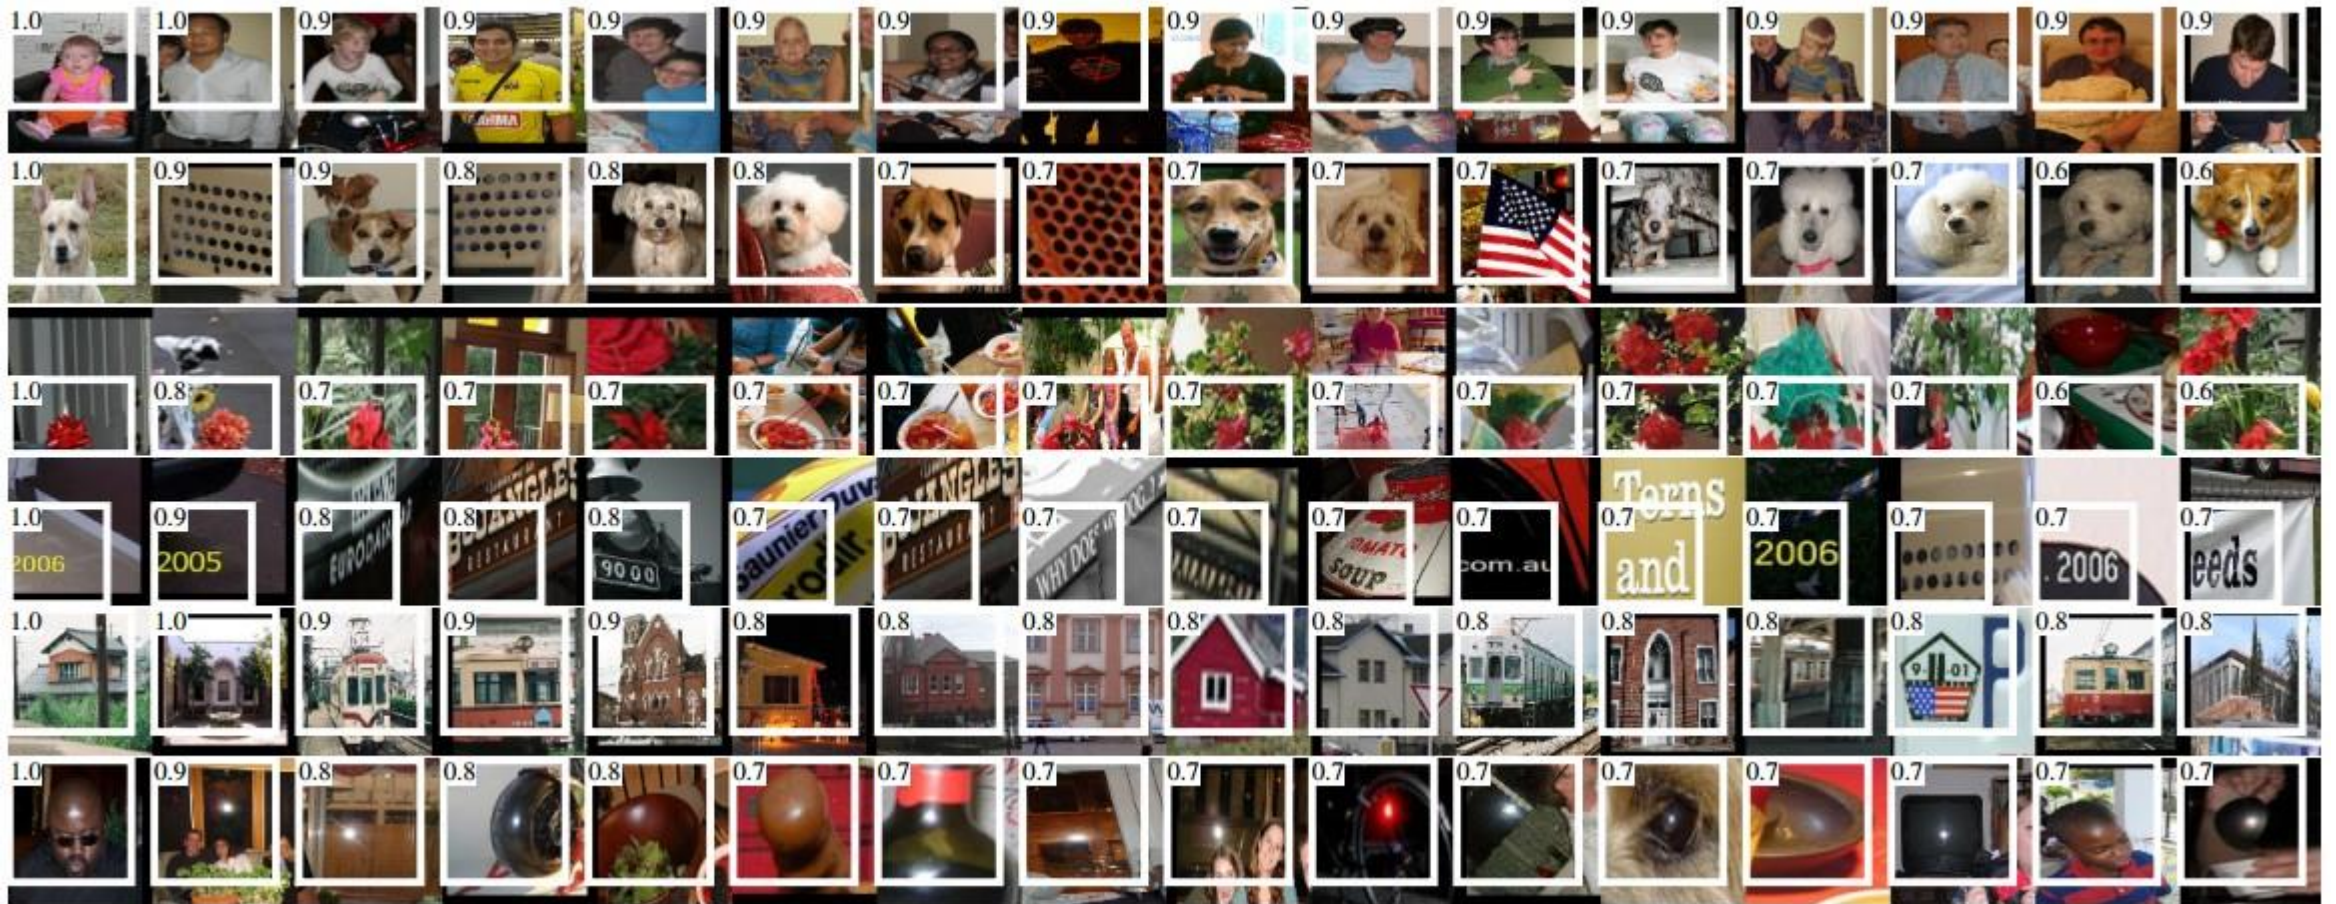
Typical-looking filters on the
trained first layer

预训练模型 首层特征图

11 x 11
(AlexNet)

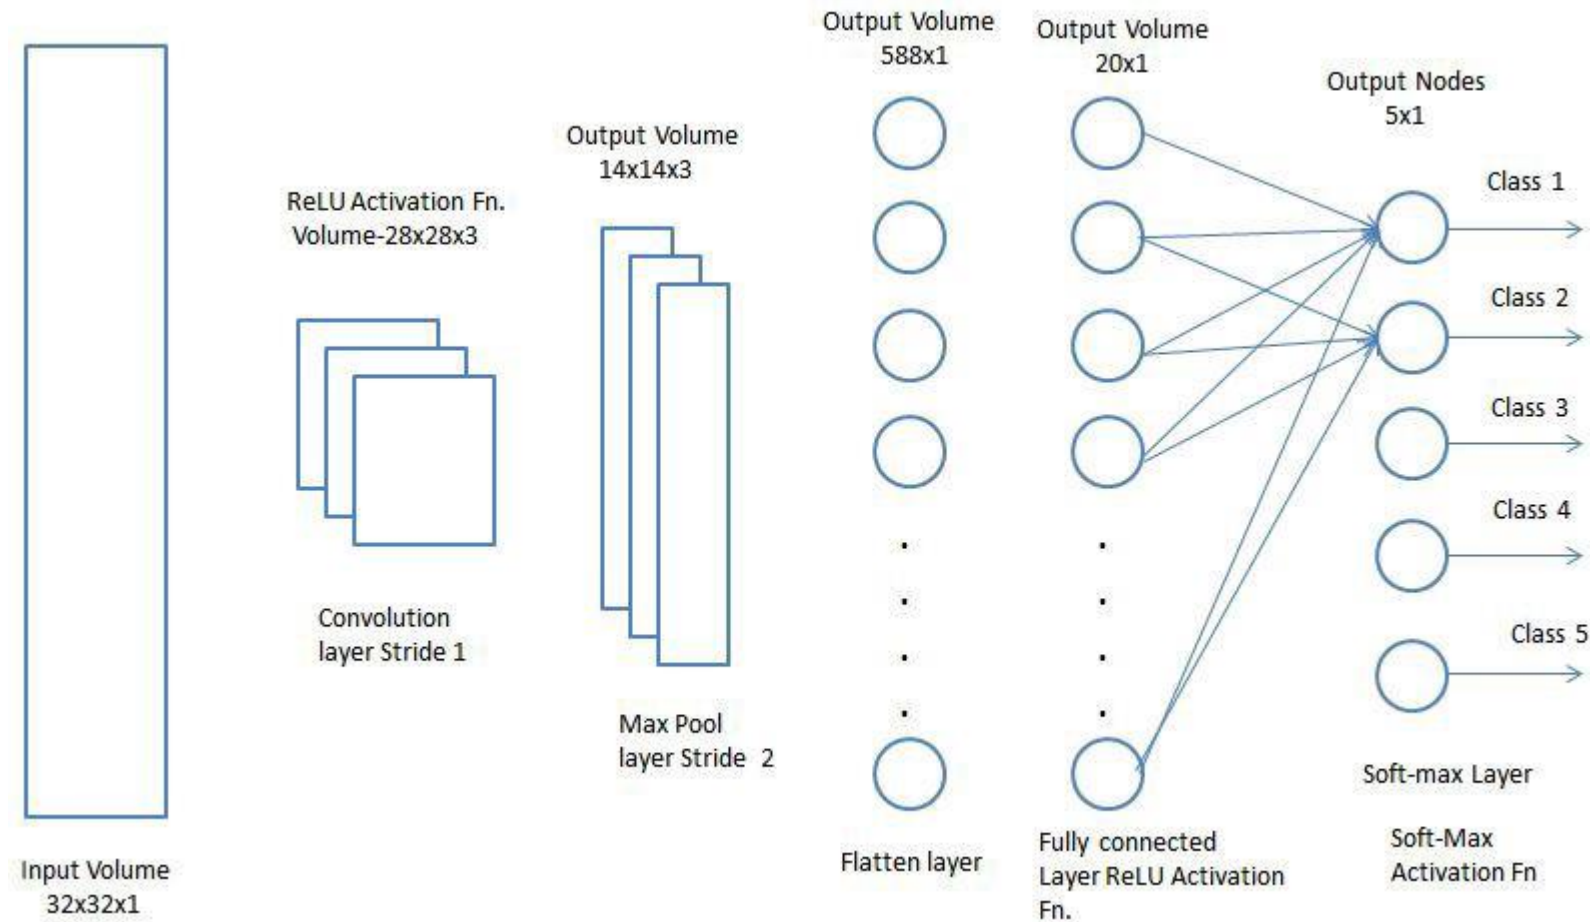


Features



Maximally activating images for some POOL5 (5th pool layer) neurons of an AlexNet.

Fully Connected Layer (FC)




More Applications: Deep Style



+



=



<https://deepdreamgenerator.com/>

More Applications: Deep Dream



+

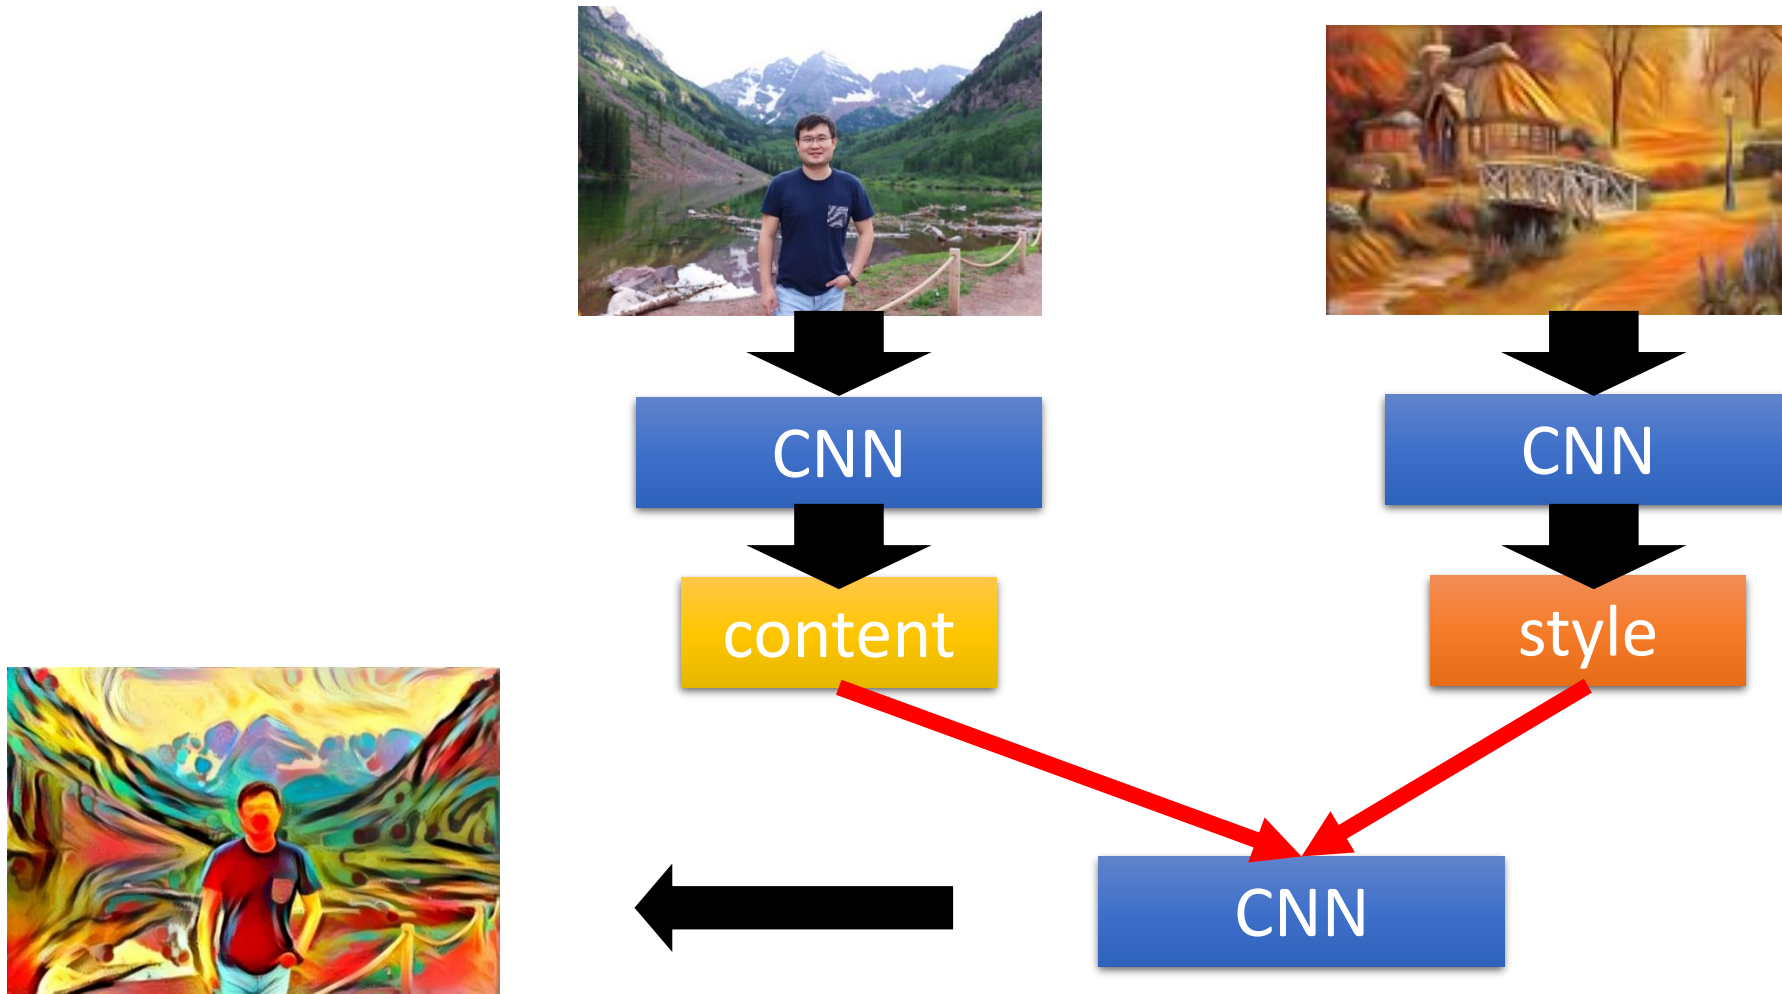


=



<https://deepdreamgenerator.com/>


Given a photo, make its style like famous paintings



Deep Learning for Computer Vision




Object Recognition



Object Detection

Deep Learning for Computer Vision



Face Detection




Skelton Detection

Deep Learning for Computer Vision



Video Generation




Unmanned Vehicle

Deep Learning for Computer Vision



Spring 2023

Deep Learning for Computer Vision




Pixel Recursive Super Resolution



Image Translation

Deep Learning for Computer Vision



Generative Adversarial Networks



Cycle GAN



02

Generative Adversarial Networks

生成对抗网络

Spring 2023

Generative Adversarial Networks



Game development and animation production are expensive and **hire many** production **artists** for relatively routine tasks.

GAN can **auto-generate** and **colorize** Anime characters. GAN可以自动生成彩色卡通人物。



(a) MidiNet model 1



(b) MidiNet model 2



(c) MidiNet model 3

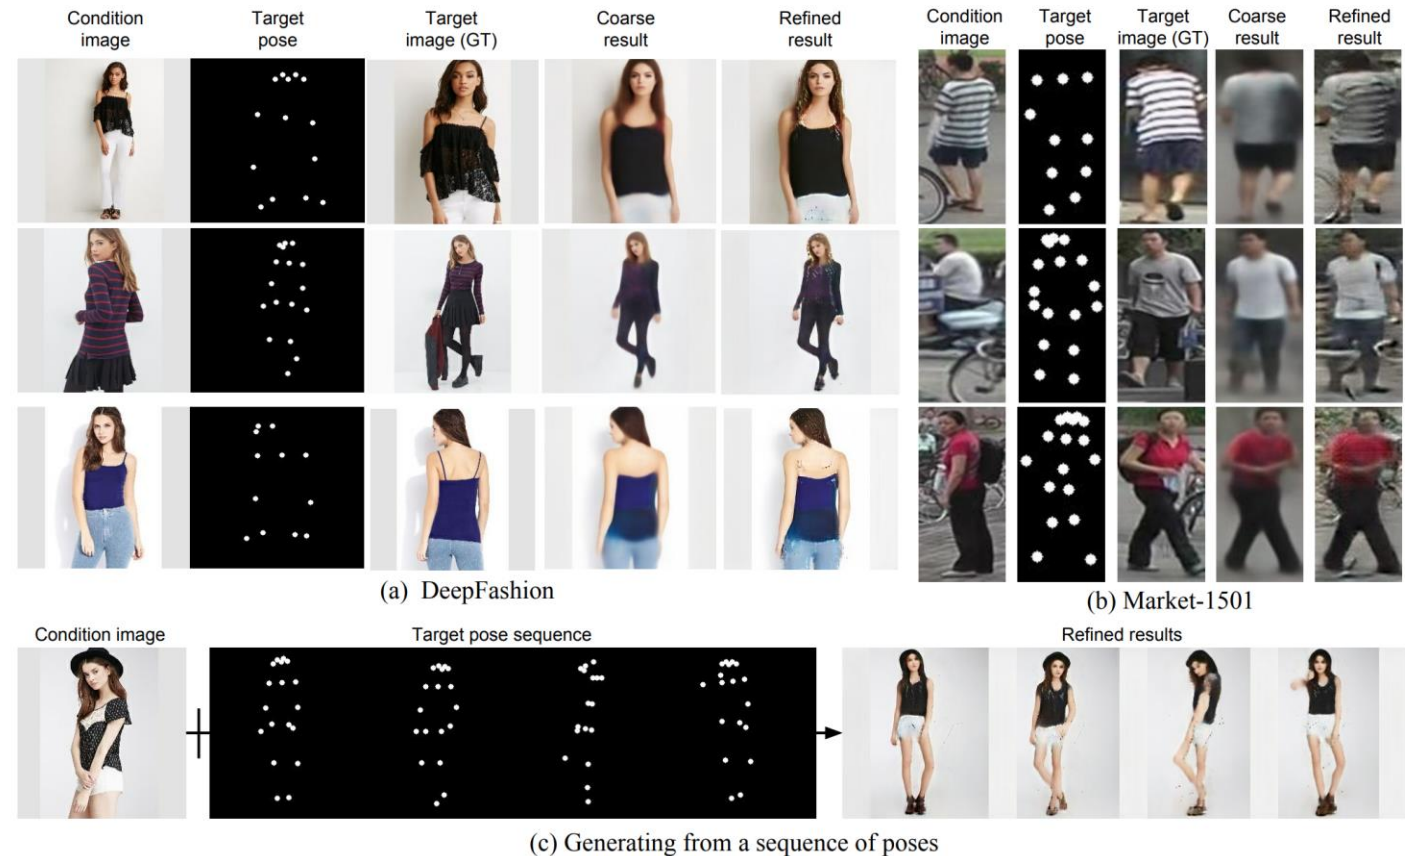


Pose Guided Person Image Generation



With an additional input of the pose, we can **transform** an image into different **poses**.

输入一个姿势，我们可以将图像中的人物变换出不同的姿势。



CycleGAN



Cross-domain transfer GANs will be likely the first batch of commercial applications. These GANs **transform images** from one domain (say real scenery) to another domain (Monet paintings or Van Gogh).

跨场景迁移图像风格



Photograph



Monet



Van Gogh

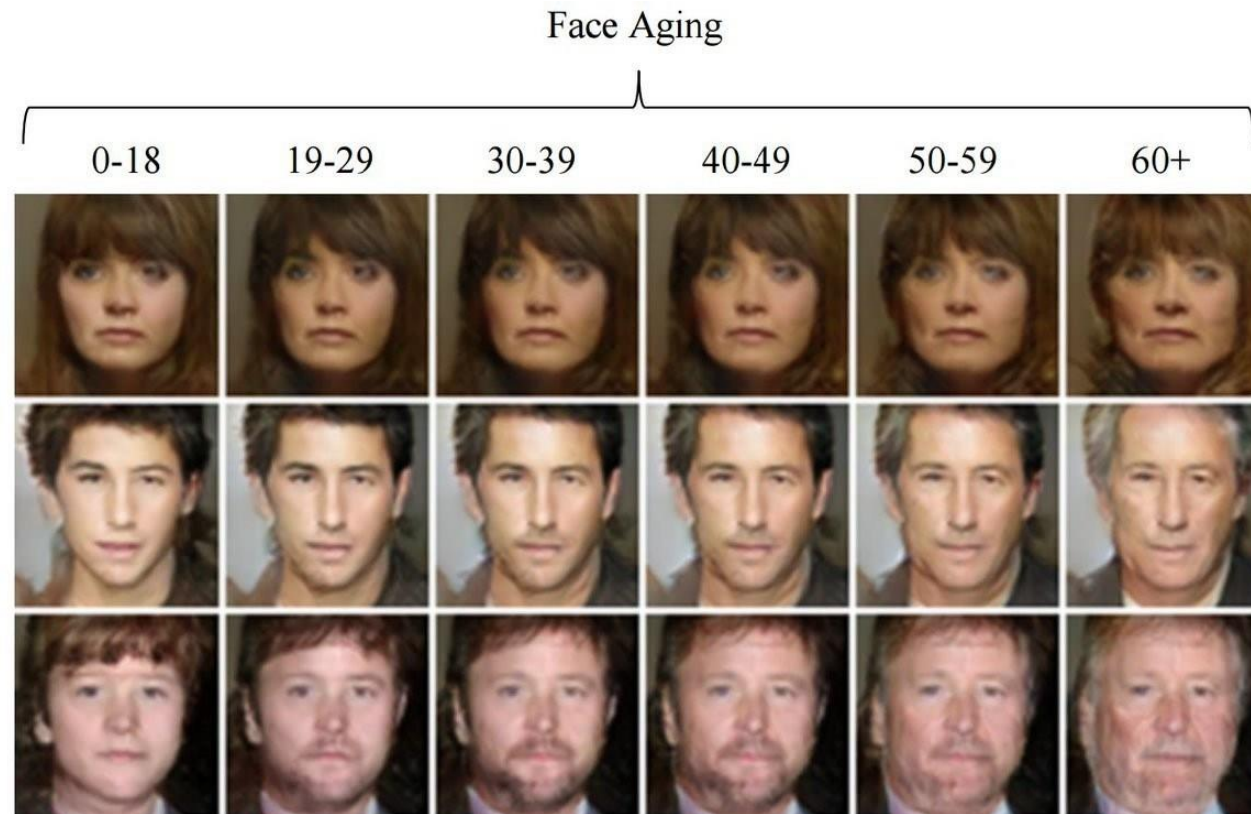


Cezanne



Ukiyo-e

Cycle GAN



PixelDTGAN



Suggesting merchandise based on celebrity pictures has been popular for fashion blogger and e-commerce.
给电商生成不同风格的衣服图片。

PixelDTGAN **creates clothing images and styles** from an image. 从一张图像上生成不同风格的衣服图像



A source image.



Possible target images.

StackGAN



Text to image is one of the earlier application of domain-transfer GAN. We input a sentence and generate multiple images fitting the description.

文字（输入）生成图像（输出）。

This bird is black with green and has a very short beak

Stage-I
images



Stage-II
images



DiscoGAN



DiscoGAN provides **matching style**: many potential applications. 风格配对

DiscoGAN learns cross domain relationship without labels or pairing. For example, it successfully transfers style (or patterns) from one domain (handbag) to another (shoe). 无需人工标注



(b) Handbag images (input) & **Generated** shoe images (output)

Generative Adversarial Networks



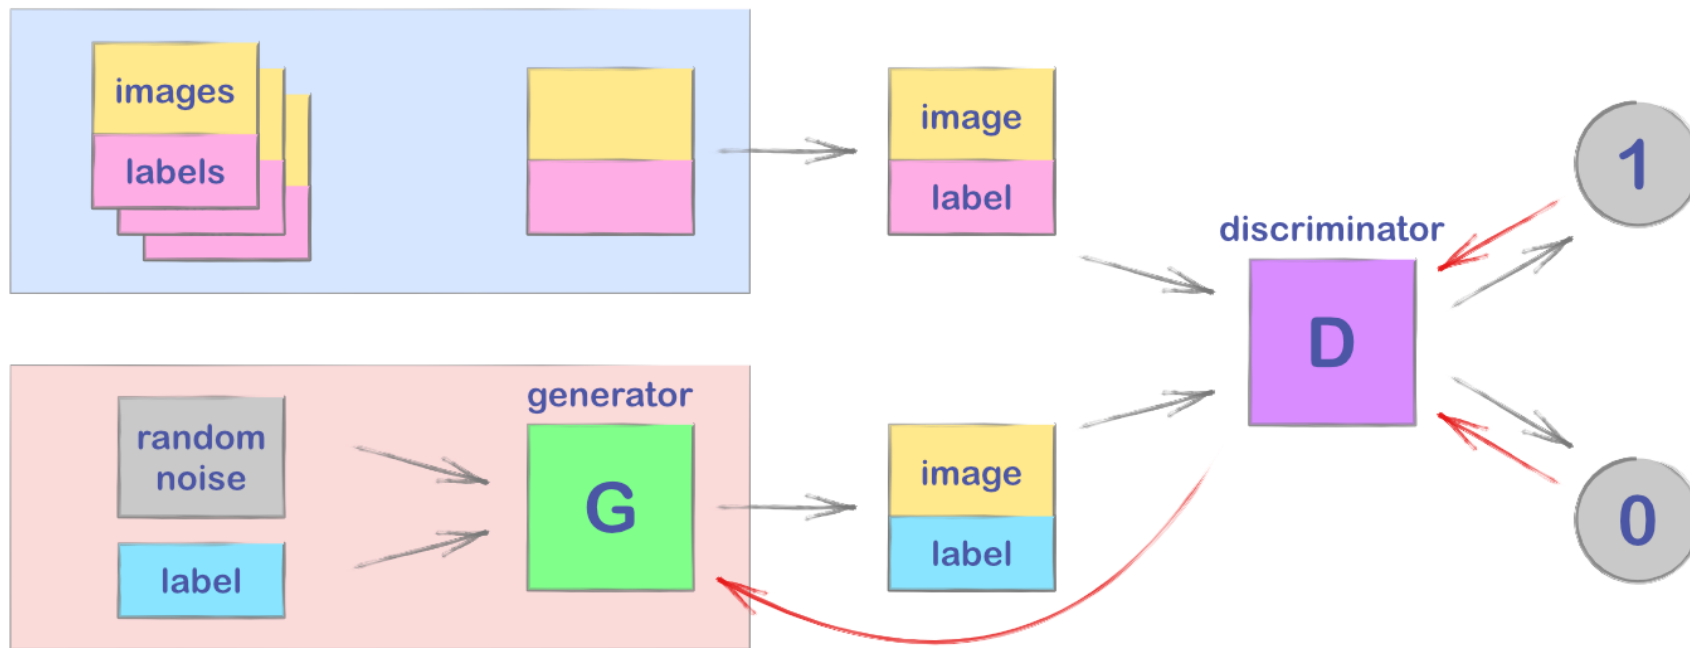
GANs have various applications on commercial market. 对抗生成网络在商业市场中无处不在。

Welcome to GANs' world. 欢迎来到对抗生成的世界。



Figure 1: Class-conditional samples generated by our model.

Basic Ideas of GAN




Generator




Image Generation 图像生成

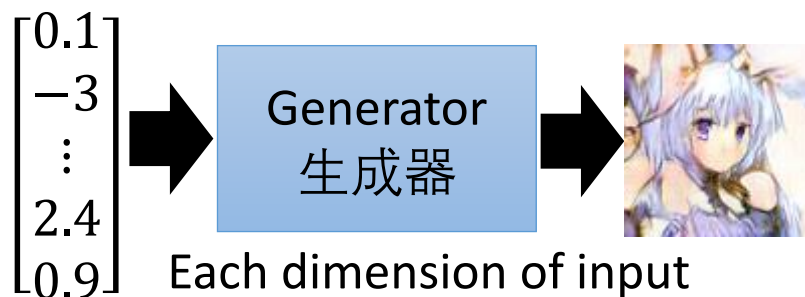
We will control what to generate latter. →
Conditional Generation 条件生成



Sentence Generation 句子生成






Generator



Each dimension of input vector represents some characteristics.



Discriminator

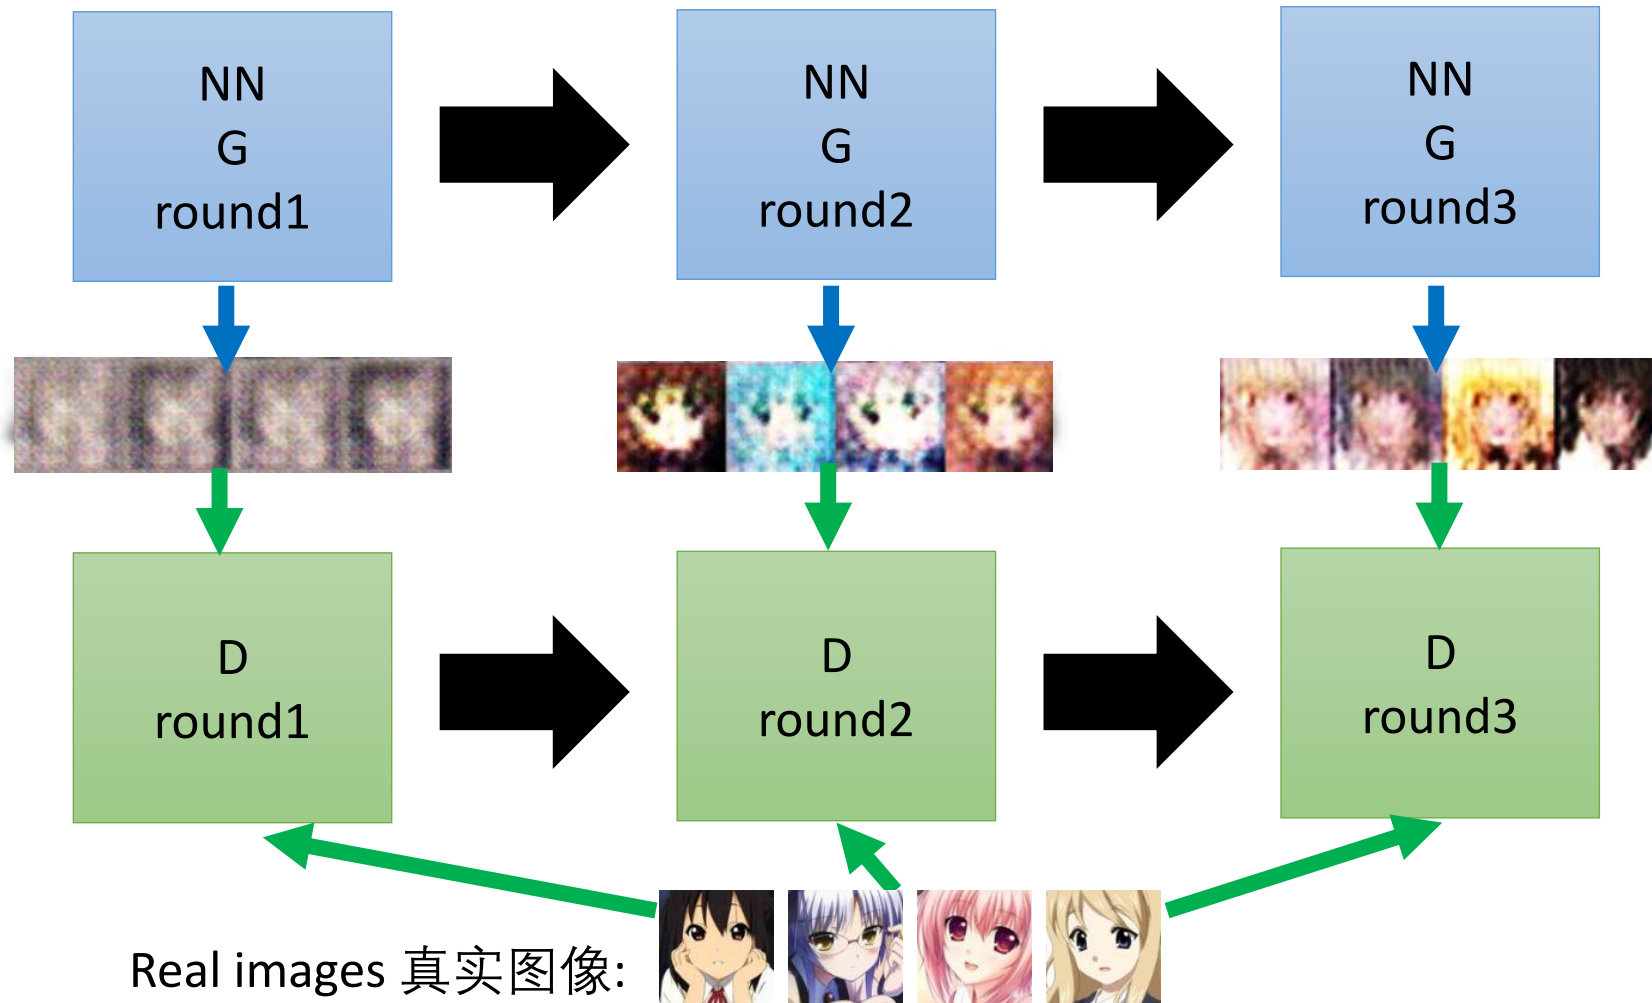


Generator vs Discriminator





Spring 2023

DiscoGAN

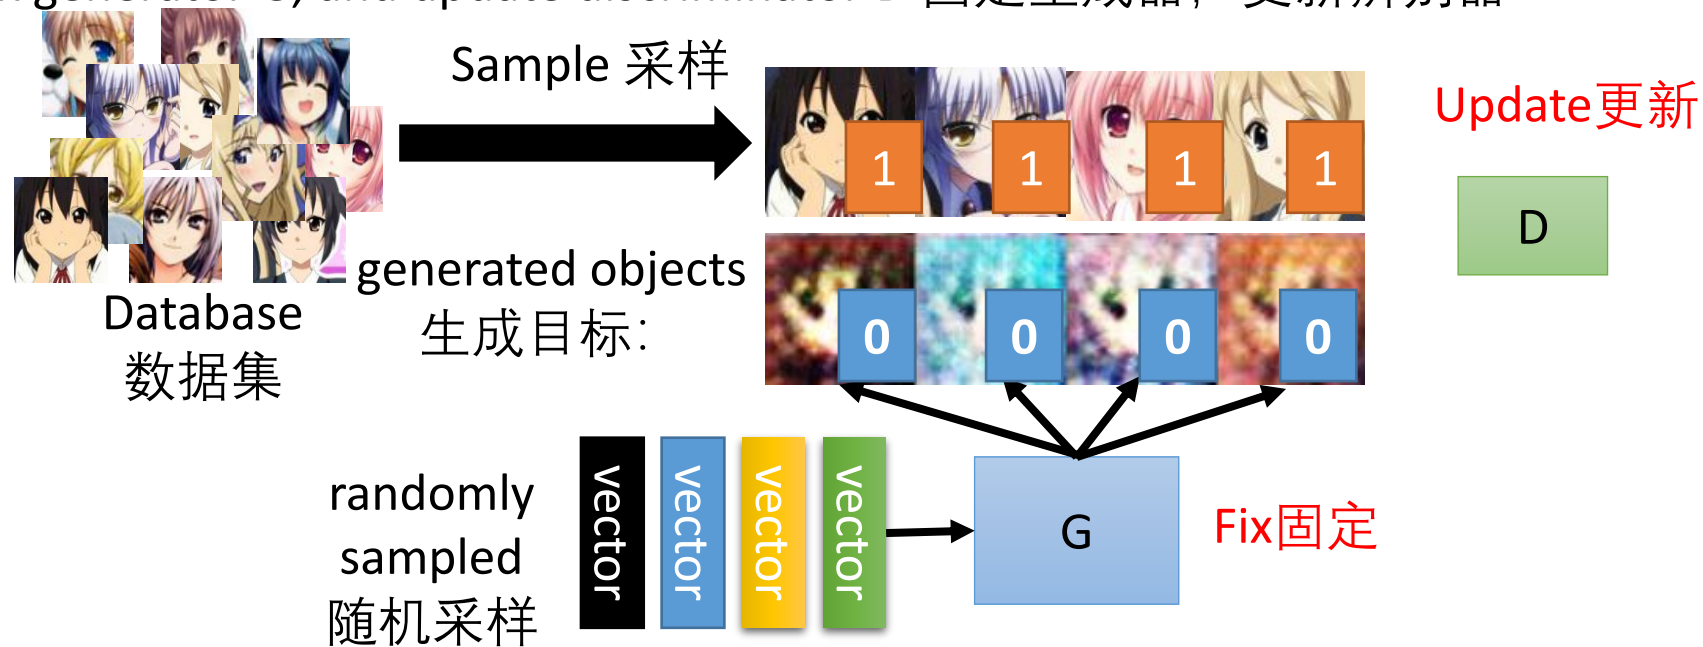


Algorithm



- Initialize generator and discriminator 初始化:  
- In each training iteration 对于训练中的每次迭代:

Step 1: Fix generator G, and update discriminator D 固定生成器, 更新辨别器



Discriminator learns to assign high scores to real objects and low scores to generated objects. 辨别器学习到如何对真实目标打高分, 对生成目标打低分。

Algorithm



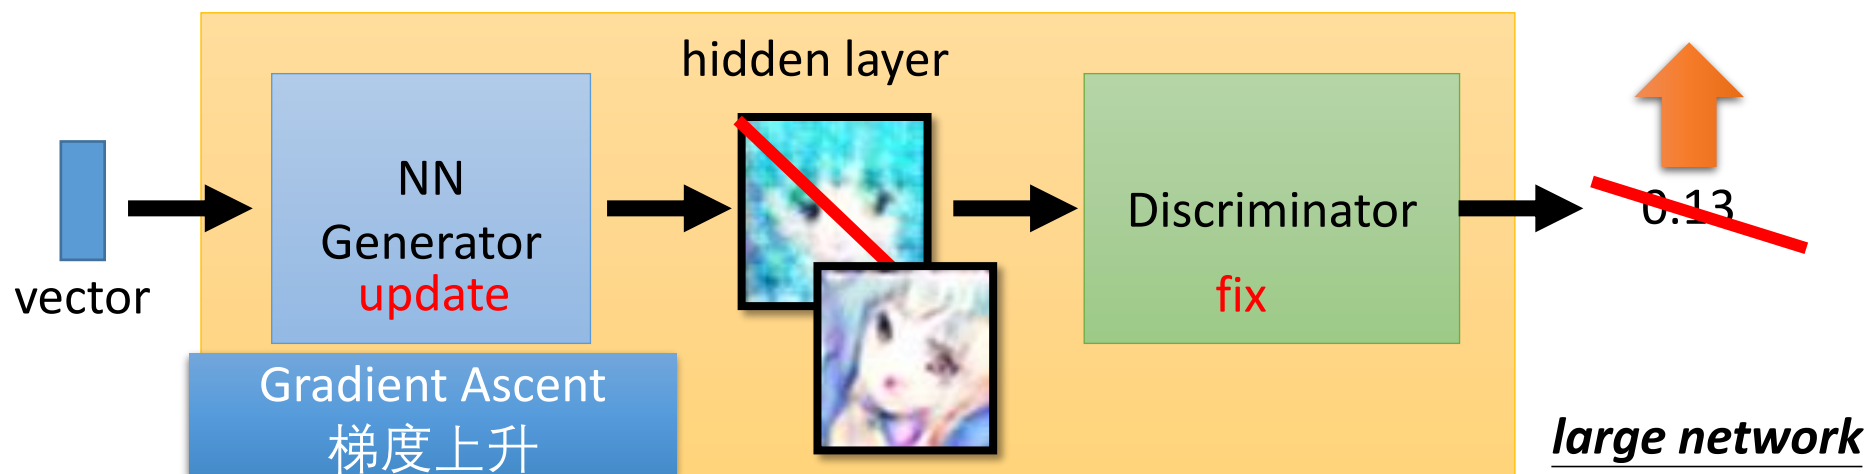
- Initialize generator and discriminator



- In each training iteration:

Step 2: Fix discriminator D, and update generator G 固定辨别器，更新生成器。

Generator learns to “fool” the discriminator 生成器学习如何去耍弄辨别器



Algorithm



Initialize θ_d for D and θ_g for G

- In each training iteration:


Learning
D

- Sample m examples $\{x^1, x^2, \dots, x^m\}$ from database
- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
- Obtaining generated data $\{\tilde{x}^1, \tilde{x}^2, \dots, \tilde{x}^m\}$, $\tilde{x}^i = G(z^i)$
- Update discriminator parameters θ_d to maximize
 - $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log D(x^i) + \frac{1}{m} \sum_{i=1}^m \log (1 - D(\tilde{x}^i))$
 - $\theta_d \leftarrow \theta_d + \eta \nabla \tilde{V}(\theta_d)$
- Sample m noise samples $\{z^1, z^2, \dots, z^m\}$ from a distribution
- Update generator parameters θ_g to maximize


Learning
G

- $\tilde{V} = \frac{1}{m} \sum_{i=1}^m \log (D(G(z^i)))$
- $\theta_g \leftarrow \theta_g - \eta \nabla \tilde{V}(\theta_g)$

GAN



Generative Adversarial Networks



Spring 2023

Q&A



Spring 2023